

Applied Statistics in R

Elena Parilina

Master's Program

Game Theory and Operations Research

Saint Petersburg State University

2020

Agenda

- ① One-Sample Tests for Normality
- ② One-Sample Tests for Exponentiality
- ③ One-Sample Tests for Discrete Probability Distributions
- ④ Two-Sample Tests
- ⑤ Tests Related to Distribution Parameters

One-Sample Tests for Normality

One-sample Kolmogorov–Smirnov test

$$D^+ = \max_{i=1,\dots,n} \left[\frac{i}{n} - F(x_{(i)}) \right], \quad D^- = \max_{i=1,\dots,n} \left[F(x_{(i)}) - \frac{i-1}{n} \right],$$

$$D = \max\{D^+, D^-\}.$$

Test statistics: $D\sqrt{n}$.

`ks.test`

```
ks.test(x, y, ...,
        alternative = c("two.sided", "less", "greater"),
        exact = NULL)
```

ks.test

- Distribution is continuous.
- Missing values are silently omitted.
- The possible values "two.sided", "less" and "greater" of alternative specify the null hypothesis that the true distribution function of x is equal to, not less than or not greater than the hypothesized distribution function.
- If `exact = NULL` (the default), an exact p-value is computed if the sample size is less than 100 in the one-sample case.

Examples

```
x <- rnorm(50)
# Does x come from a shifted gamma distribution with shape
# 3 and rate 2?
ks.test(x+2, "pgamma", 3, 2) # two-sided, exact
ks.test(x+2, "pgamma", 3, 2, exact = FALSE)
ks.test(x+2, "pgamma", 3, 2, alternative = "gr")
```

Lilliefors (one-sample Kolmogorov–Smirnov) test

$$D^+ = \max_{i=1,\dots,n} \left[\frac{i}{n} - F_0\left(\frac{x_{(i)} - \bar{x}}{s}\right) \right], \quad D^- = \max_{i=1,\dots,n} \left[F_0\left(\frac{x_{(i)} - \bar{x}}{s}\right) - \frac{i-1}{n} \right],$$

$$D = \max\{D^+, D^-\}.$$

Test statistics: $D \left(\sqrt{n} - 0.01 + \frac{0.85}{\sqrt{n}} \right).$

lillie.test

```
library(nortest)
lillie.test(x)
```

uniNorm

```
library(MVN)
uniNorm(data, type = "Lillie", desc = TRUE)
```

`lillie.test`

- Argument: a numeric vector of data values, the number of which must be greater than 4. Missing values are allowed.
- Although the test statistic obtained from `lillie.test(x)` is the same as that obtained from `ks.test(x, "pnorm", mean(x), sd(x))`, it is not correct to use the p -value from the latter for the composite hypothesis of normality (mean and variance unknown), since the distribution of the test statistic is different when the parameters are estimated.

Examples

```
lillie.test(rnorm(100, mean = 5, sd = 3))  
lillie.test(runif(100, min = 2, max = 4))
```

uniNorm

- This function performs univariate normality tests, including Shapiro-Wilk, Cramer-von Mises, Lilliefors (Kolmogorov-Smirnov), Shapiro-Francia and Anderson-Darling.
- Arguments: data is a vector, data frame or matrix; type select one of the univariate normality tests: SW Shapiro-Wilk, CVM Cramer-von Mises, Lillie Lilliefors (Kolmogorov-Smirnov), SF Shapiro-Francia, AD Anderson-Darling; desc if TRUE, it displays descriptive statistics including mean, standard deviation, median, minimum, maximum, 25th and 75th percentiles, skewness and kurtosis.

```
uniNorm(x, type = "SW", desc = TRUE)
```


Anderson–Darling test

$$A = -n - \frac{1}{n} \sum_{i=1}^n (2i-1) \left[\ln F_0 \left(\frac{x_{(i)} - \bar{x}}{s} \right) + \ln(1 - F_0 \left(\frac{x_{(n-i+1)} - \bar{x}}{s} \right)) \right]$$

Test statistics: $A \left(1 + \frac{0.75}{n} + \frac{2.25}{n^2} \right)$.

ad.test

```
library(nortest)
ad.test(x)
```

uniNorm

```
library(MVN)
uniNorm(data, type = "AD", desc = TRUE)
```

ad.test

- Argument: `x` which is a numeric vector of data values, the number of which must be greater than 7. Missing values are allowed.
- Compared to the Cramer-von Mises test (as second choice) it gives more weight to the tails of the distribution.

Examples

```
ad.test(rnorm(100, mean = 5, sd = 3))
```

```
ad.test(runif(100, min = 2, max = 4))
```

Cramer–von Mises test

$$\omega^2 = \frac{1}{12n} + \sum_{i=1}^n \left(F_0 \left(\frac{x_{(i)} - \bar{x}}{s} \right) - \frac{2i-1}{2n} \right)^2.$$

Test statistics: $\omega^2 \left(1 + \frac{0.5}{n} \right)$.

cvm.test

```
library(nortest)
  cvm.test(x)
cvm.test(rnorm(100, mean = 5, sd = 3))
cvm.test(runif(100, min = 2, max = 4))
```

uniNorm

```
library(MVN)
  uniNorm(data, type = "CVM", desc = TRUE)
```

Pearson chi-square test

Test statistics:

$$\sum_{i=1}^r \frac{\left(n_i - np_i^{(0)}(\hat{\theta})\right)^2}{np_i^{(0)}(\hat{\theta})}.$$

`pearson.test`

```
library(nortest)
pearson.test(x, n.classes = ceiling(2 * (n2/5)),
             adjust = TRUE)
```

- `n.classes` is a number of classes. The default is due to Moore (1986).
- `adjust` is logical; if `TRUE` (default), the p -value is computed from a chi-square distribution with `n.classes-3` degrees of freedom, otherwise from a chi-square distribution with `n.classes-1` degrees of freedom.

pearson.test

The Pearson chi-square test is usually not recommended for testing the composite hypothesis of normality due to its inferior power properties compared to other tests. It is common practice to compute the p-value from the chi-square distribution with $n.\text{classes} - 3$ degrees of freedom, in order to adjust for the additional estimation of two parameters. (For the simple hypothesis of normality (mean and variance known) the test statistic is asymptotically chi-square distributed with $n.\text{classes} - 1$ degrees of freedom.) This is, however, not correct as long as the parameters are estimated by $\text{mean}(x)$ and $\text{var}(x)$ (or $\text{sd}(x)$), as it is usually done, see Moore (1986) for details.

Since the true p-value is somewhere between the two, it is suggested to run `pearson.test` twice, with `adjust = TRUE` (default) and with `adjust = FALSE`. It is also suggested to slightly change the default number of classes, in order to see the effect on the p-value. Eventually, it is suggested not to rely upon the result of the test.

Shapiro–Wilk test:

shapiro.test

```
shapiro.test(x)
```

uniNorm

```
library(MVN)  
uniNorm(data, type = "SW", desc = TRUE)
```

Shapiro–Francia test:

sf.test

```
library(nortest)  
sf.test(x)
```

uniNorm

```
library(MVN)  
uniNorm(data, type = "SF", desc = TRUE)
```

One-Sample Tests for Exponentiality

One-sample Kolmogorov–Smirnov test

Test statistics:

$$KS = \sup_{x>0} |F_n(x) - (1 - e^{-x})|,$$

where F_n is the ecdf of the scaled data $y_i = \frac{x_i}{\bar{x}}$.

The test uses Monte Carlo simulations.

`ks.exp.test`

```
library(exptest)
ks.exp.test(x, nrepl=2000)
ks.exp.test(rexp(100))
ks.exp.test(runif(100, min = 50, max = 100))
```

`nrepl` is the number of replications in Monte Carlo simulation.

Cramer–von Mises test

$$\omega^2 = \int_0^\infty [F_n(x) - (1 - e^{-x})]^2 e^{-x} dx,$$

where F_n is the ecdf of the scaled data $y_i = \frac{x_i}{\bar{x}}$.

The test uses Monte Carlo simulations.

`cvm.exp.test`

```
library(exptest)
cvm.exp.test(x, nrepl=2000)
cvm.exp.test(rexp(100))
cvm.exp.test(runif(100, min = 50, max = 100))
```

`nrepl` is the number of replications in Monte Carlo simulation.

Atkinson test test

Test statistics:

$$T(p) = \sqrt{n} \left| \frac{\left(\frac{1}{n} \sum_{i=1}^n x_i^p \right)^{1/p}}{\bar{x}} - (\Gamma(1+p))^{1/p} \right|$$

The test uses Monte Carlo simulations.

`atkinson.exp.test`

```
library(exptest)
atkinson.exp.test(x, p=0.99, simulate.p.value=FALSE,
nrepl=2000)
atkinson.exp.test(rexp(100))
atkinson.exp.test(rchisq(100,3))
```

`simulate.p.value` is a logical value indicating whether to compute p -values by Monte Carlo simulation.

Lorenz test

$$L = \frac{\sum_{i=1}^{np} x_{(i)}}{\sum_{i=1}^n x_{(i)}}.$$

Test statistics: $\sqrt{n}(L - p - (1 - p) \ln(1 - p))$.

The test uses Monte Carlo simulations.

`lorenz.exp.test`

```
library(exptest)
lorenz.exp.test(x, p=0.5, simulate.p.value=FALSE,
  nrepl=2000)
```

Shapiro-Wilk test for exponentiality:

`shapiro.exp.test`

```
library(exptest)
shapiro.exp.test(x, nrepl=2000)
```

Kimber-Michael test for exponentiality:

`shapiro.exp.test`

```
library(exptest)
kimber.exp.test(x, nrepl=2000)
```

One-Sample Tests for Discrete Probability Distributions

Pearson chi-square test

goodfit

```
library(vcd)
gf<-goodfit(x,type="poisson", "MinChisq")
summary(gf)

# type = c("poisson", "binomial", "nbinomial")
# method = c("ML", "MinChisq")
```

The negative binomial distribution with size= n and prob= p has density

$$\frac{\Gamma(x+n)}{\Gamma(n)x!} p^n (1-p)^x$$

for $x = 0, 1, 2, \dots, n > 0$ and $0 < p \leq 1$.

This represents the number of failures which occur in a sequence of Bernoulli trials before a target number of successes is reached. The mean is $\mu = n(1-p)/p$ and variance $n(1-p)/p^2$. This definition allows non-integer values of size

goodfit

Examples

```
# Simulated data examples:
dummy <- rnbinoom(200, size = 5, prob = 0.8)
gf <- goodfit(dummy, type = "nbinomial", method =
"MinChisq")

dummy <- rbinom(100, size = 6, prob = 0.5)
gf1 <- goodfit(dummy, type = "binomial", par = list(size =
6))
gf2 <- goodfit(dummy, type = "binomial", par = list(prob =
0.6, size = 6))
```

Two-Sample Tests

Two-sample Kolmogorov–Smirnov test

$$D^+ = \max_{r=1,\dots,m} \left[\frac{r}{m} - F_n(y_{(r)}) \right] = \max_{s=1,\dots,n} \left[G_m(x_{(s)}) - \frac{s-1}{n} \right],$$

$$D^- = \max_{r=1,\dots,m} \left[F_n(y_{(r)}) - \frac{r-1}{m} \right] = \max_{s=1,\dots,n} \left[\frac{s}{n} - G_m(x_{(s)}) \right],$$

$$D = \max\{D^+, D^-\}.$$

Test statistics: $D\sqrt{\frac{mn}{m+n}}.$

`ks.test`

```
ks.test(x, y, ...,
        alternative = c("two.sided", "less", "greater"),
        exact = NULL)
```

Wilcoxon rank sum test / Mann–Whitney U-test

`ks.test`

```
wilcox.test(x, y = NULL,  
            alternative = c("two.sided", "less", "greater"),  
            mu = 0, paired = FALSE, exact = NULL, correct = TRUE,  
            conf.int = FALSE, conf.level = 0.95, ...)
```

Tests Related to Distribution Parameters

F-test

Test statistics:

$$F = \frac{s_1^2 n / (n - 1)}{s_2^2 m / (m - 1)} = \frac{\tilde{s}_1^2}{\tilde{s}_2^2}.$$

var.test

```
var.test(x, y, ratio = 1,  
         alternative = c("two.sided", "less", "greater"),  
         conf.level = 0.95, ...)
```

Bartlett's test

This is an analog of F-test for more than two samples.

`bartlett.test`

```
bartlett.test(x, g, ...)
```

or

```
bartlett.test(formula, data, subset, ...)
```

`bartlett.test`

- `x` — a numeric vector of data values, or a list of numeric data vectors representing the respective samples, or fitted linear model objects (inheriting from class "lm").
- `g` — a vector or factor object giving the group for the corresponding elements of `x`. Ignored if `x` is a list.
- `formula` — a formula of the form `lhs rhs` where `lhs` gives the data values and `rhs` the corresponding groups.
- `data` — an optional matrix or data frame (or similar: see

Student's t-test

One-sample t-test: $\tau = \frac{\bar{x} - a_0}{s} \sqrt{n-1} = \frac{\bar{x} - a_0}{\tilde{s}} \sqrt{n}$

t.test

```
t.test(x, alternative = c("two.sided", "less", "greater"),
      mu = 0, conf.level = 0.95, ...)
```

Two-sample t-test:

$$\tau_1 = \frac{\bar{x} - \bar{y}}{\sqrt{\sigma_1^2/n + \sigma_2^2/m}}, \quad \tau_2 = \frac{\bar{x} - \bar{y}}{\hat{s} \sqrt{1/n + 1/m}},$$

$$\hat{s} = \sqrt{\frac{ns_1^2 + ms_2^2}{n+m-2}} = \sqrt{\frac{(n-1)\tilde{s}_1^2 + (m-1)\tilde{s}_2^2}{n+m-2}}.$$

t.test

```
t.test(x, y = NULL, alternative = c("two.sided", "less",
  "greater"), mu = 0, paired = FALSE, var.equal = FALSE,
  conf.level = 0.95, ...)
```

Pairwise t-test

`pairwise.t.test`

```
pairwise.t.test(x, g, p.adjust.method = p.adjust.methods,
pool.sd = !paired, paired = FALSE, alternative =
c("two.sided", "less", "greater"), ...)
```

To test two-sample t-test for pairs and to have a resulting table of p-val.

`pairwise.t.test`

- `x` — response vector.
- `g` — grouping vector or factor.
- `p.adjust.method` — Method for adjusting p values (see `p.adjust`).
- `pool.sd` — switch to allow/disallow the use of a pooled SD
- `paired` — a logical indicating whether you want paired t-tests.
- `alternative` — a character string specifying the alternative hypothesis, must be one of "two.sided" (default), "greater" or "less".

Probability of success in a Bernoulli experiment

Test statistics: $B = \sum_{i=1}^n x_i.$

`binom.test`

```
binom.test(x, n, p = 0.5,  
  alternative = c("two.sided", "less", "greater"),  
  conf.level = 0.95)
```


Sign Test for the population median

Test statistics:

$$ST = \sum_{i=1}^n s(x_i - \theta_0),$$

where

$$s(x_i - \theta_0) = \begin{cases} 1, & x_i > \theta_0, \\ 0, & x_i \leq \theta_0. \end{cases}$$

binom.test

```
library(BSDA) SIGN.test(x, md = 0,  
  alternative = c("two.sided", "less", "greater"),  
  conf.level = 0.95)
```